

An Intro to DifferentialEquations.jl

Chris Rackauckas

June 14, 2020

0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable u changes, that is

$$u' = f(u, p, t)$$

where p are the parameters of the model, t is the time variable, and f is the nonlinear model of how u changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value $u(0) = u_0$. Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now $t = 0$ and measuring time in years, our model is:

$$u' = 0.98u$$

and $u(0) = 1.0$. We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.

```
sol = solve(prob)
```


```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.100424944449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

Analyzing the Solution Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the

functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:

```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in  $\mu\text{m}$ )",label="My Thick Line!") # legend=false
```


Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is $u(t) = u_0 \exp(at)$, so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```

Solution to the linear ODE with a thick line



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t
```

```
5-element Array{Float64,1}:  
 0.0  
 0.10042494449239292  
 0.35218603951893646  
 0.6934436028208104  
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u
```

```
5-element Array{Float64,1}:  
 1.0  
 1.1034222047865465  
 1.4121908848175448  
 1.9730384275622996  
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:
 1.0
 1.2038471492789395
 1.7643769243364813
 2.66648203038311
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)
```

```
1.554261048055312
```

Controlling the Solver `DifferentialEquations.jl` has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances `abstol` and `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally, `reltol` is the relative accuracy while `abstol` is the accuracy when `u` is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults


`abstol=1e-6` and `reltol=1e-3`, you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob,abstol=1e-8,reltol=1e-8)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
 0.28631546412766684
 0.4381941361169628
 0.6118924302028597
 0.7985659100883337
 0.9993516479536952
 1.0
u: 9-element Array{Float64,1}:
 1.0
 1.0412786454705882
 1.1547261252949712
 1.3239095703537043
 1.5363819257509728
 1.8214895157178692
 2.1871396448296223
 2.662763824115295
 2.664456241933517
```

Now we can see no visible difference against the true solution:

```
plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of $t=0.1k$ for integers k , we would use the command:

```
sol = solve(prob,saveat=0.1)
```

```
retcode: Success
Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
```



```
1.4799379510586077
1.632316207054161
1.8003833264983584
1.9857565541588758
2.1902158127997695
2.415725742084496
2.664456142481451
```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```
retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496
```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```
sol = solve(prob,dense=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```

retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451

```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```

retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451

```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

Choosing Solver Algorithms There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian `f` with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in `f` (like parameters of order `1e5`), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```

retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142

```

```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```
sol = solve(prob,Tsit5(),reltol=1e-6)
```

```

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition $u_0 = [1.0, 0.0, 0.0]$ as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters $\sigma = 10$, $\rho = 28$, and $\beta = 8/3$, and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!
```

```
(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```

ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]

```

Now, just as before, we solve the problem:

```
sol = solve(prob)
```

```

retcode: Success
Interpolation: Automatic order switching interpolation
t: 1294-element Array{Float64,1}:
 0.0
 3.5678604836301404e-5
 0.0003924646531993154
 0.0032624077544510573
 0.009058075635317072
 0.01695646895607931
 0.0276899566248403
 0.041856345938267966
 0.06024040228733675
 0.08368539694547242
 ⋮
99.39403070915297
99.47001147494375
99.54379656909015
99.614651558349
99.69093823148101
99.78733023233721
99.86114450046736
99.96115759510786
100.0
u: 1294-element Array{Array{Float64,1},1}:
 [1.0, 0.0, 0.0]
 [0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
 [0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
 [0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
 [0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
 [0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
 [0.8483309877783048, 0.69156288756671, 0.008487623500490047]
 [0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
 [0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
 [1.0888636764765296, 2.052326153029042, 0.07402570506414284]
 ⋮
 [12.999157033749652, 14.10699925404482, 31.74244844521858]
 [11.646131422021162, 7.2855792145502845, 35.365000488215486]
 [7.777555445486692, 2.5166095828739574, 32.030953593541675]
 [4.739741627223412, 1.5919220588229062, 27.249779003951755]
 [3.2351668945618774, 2.3121727966182695, 22.724936101772805]
 [3.310411964698304, 4.28106626744641, 18.435441144016366]
 [4.527117863517627, 6.895878639772805, 16.58544600757436]
 [8.043672261487556, 12.711555298531689, 18.12537420595938]
 [9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the *i*th time point.

```
sol.t[10],sol[10]
```

```
(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506
414284])
```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the *j*th variable at time *i*:

```
sol[2,10]
```

```
2.052326153029042
```


We can get a real matrix by performing a conversion:

```
A = Array(sol)
```

```
3×1294 Array{Float64,2}:
 1.0  0.999643    0.996105    0.969359    ...    4.52712    8.04367    9.97538
 0.0  0.000998805  0.0109654    0.0897706           6.89588   12.7116   15.1439
 0.0  1.78143e-8   2.14696e-6   0.000143802       16.5854   18.1254   21.0064
```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol, vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is $u[1]$, the y axis is $u[2]$, and the z axis is $u[3]$. Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol,vars=(1,2,3),denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol,vars=(0,2))
```

0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. `DifferentialEquations.jl` respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
Interpolation: Automatic order switching interpolation
t: 10-element Array{Float64,1}:
 0.0
 0.05392648622307931
```

```

0.1406132536939904
0.23261557152239454
0.3516952255025277
0.4823733912189384
0.6275835875214623
0.796759097145997
0.9831278804494583
1.0
u: 10-element Array{Array{Float64,2},1}:
 [0.20075484072637195 0.7757405105775659; 0.8967433436433823 0.967117294956
2656; 0.980525927678076 0.3303948714318077; 0.1677649362171667 0.0703911672
6426942]
 [0.15360775609339183 0.7784140739473572; 1.006607155225334 1.0559363342164
354; 0.9532188047086694 0.16992661752982702; 0.2513257214925577 0.221682566
79909618]
 [0.02951434485930901 0.6955199499599316; 1.1040695908117368 1.058141269438
6352; 0.9463791460600056 -0.05925537968655875; 0.35210816419319 0.456759024
15339353]
 [-0.15194608228725986 0.4876549356117511; 1.1272944821492565 0.88450309624
10229; 1.0028989801660435 -0.22809030934046706; 0.40313529490278777 0.68230
23464595378]
 [-0.4212652951959783 0.042442049171868834; 1.1006611985833963 0.4418198174
0930915; 1.1860511674708485 -0.2661342597463349; 0.36664639193607607 0.9110
584902252213]
 [-0.6807811668698157 -0.640006799759824; 1.1112545192193721 -0.20922446784
223064; 1.5152771620987866 0.011554127654654689; 0.17753773665045988 1.0386
239131781112]
 [-0.7873540985438721 -1.5415177748390512; 1.3333350295151487 -0.9108926163
871542; 1.9541352687341735 0.7899997988154843; -0.21787518187205232 0.96855
29498617972]
 [-0.44798557154448604 -2.551289533555641; 2.0583729428090534 -1.2935544052
94493; 2.3165339317463496 2.318525328714581; -0.8812434028949278 0.52098350
74470793]
 [0.7756638100834994 -3.1458661250822226; 3.506235710565667 -0.577580643490
3984; 2.015026048712961 4.5048288006668695; -1.7092420295582131 -0.51004388
26147548]
 [0.9372824658462704 -3.1508581381092617; 3.6644387978403876 -0.43325426865
127825; 1.9278367149014712 4.70773601023575; -1.77977550335718 -0.632061999
471221]

```

There is no real difference from what we did before, but now in this case u_0 is a 4×2 matrix. Because of that, the solution at each time point is matrix:

```
sol[3]
```

```

4×2 Array{Float64,2}:
 0.0295143  0.69552
 1.10407    1.05814
 0.946379  -0.0592554
 0.352108  0.456759

```

In `DifferentialEquations.jl`, you can use any type that defines $+$, $-$, $*$, $/$, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change the input to be a matrix of `BigFloat`:

```
big_u0 = big.(u0)
```

```
4×2 Array{BigFloat,2}:
```

```
0.200755  0.775741
0.896743  0.967117
0.980526  0.330395
0.167765  0.0703912
```

and we can solve the `ODEProblem` with arbitrary precision numbers by using that initial condition:

```
prob = ODEProblem(f, big_u0, tspan)
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: Automatic order switching interpolation
```

```
t: 6-element Array{Float64,1}:
```

```
0.0
0.06710711556243842
0.28768651229642395
0.5831183483910013
0.9086442367054589
1.0
```

```
u: 6-element Array{Array{BigFloat,2},1}:
```

```
[0.200754840726371952541740029118955135345458984375 0.77574051057756587823
632798972539603710174560546875; 0.89674334364338226244228735595243051648139
95361328125 0.9671172949562656384614456328563392162322998046875; 0.98052592
76780760036018591563333757221698760986328125 0.3303948714318076795848355686
7577135562896728515625; 0.1677649362171667046794709676760248839855194091796
875 0.070391167264269416392608036403544247150421142578125]
```

```
[0.13836592436844159672189699497226186070842323684330617228932750996993447
98364662 0.7728229813739593359484998602699201476884690663867605232407025387
114413087713372; 1.02733484649925108098989522451727280705566022678700923500
9423167186231747808367 1.06731565429613635896144560492012739385306538904483
6171637376983944492047333065; 0.9489493766100118767572734088936885712984566
614271602238143560309675037421889613 0.132186526396001034573670718935402142
4676456196656649530111105376506820289022597; 0.2695107459447621075006758374
707234482880801259011223892499500204476265664586959 0.258287779695249366152
0639844773241310236441775887613082625417830570453813793601]
```

```
[-0.2750539748658655505189115579738245256794934097395387494526624088823932
296471318 0.305387352130809245167342621013926095377547348897870558341008787
1259543987491882; 1.1178595289049776015038874832621354529516980403034345451
97004567958347033931427 0.7057055783332081587684533095136511575496292492512
233377458583573901065372161238; 1.07208851996608955117089054874572839169835
4339514986734543733222788767238380353 -0.2753515953994391361420825452790869
341199263673577029674401735034612353360721932; 0.40160411729150074605004782
89169655596192719227404706792698527945716660557687385 0.7988449318850210955
537769851220064772643829564988609695366221070432540511816958]
```

```
[-0.7847703693366268460727037508459522108445850211271222056020074446811243
828621879 -1.25801418559606915542279920194044964863795584990495281423971492
1473916534734309; 1.2321962098943289728486648018518510538339275868524356784
03484011179062586332557 -0.715417723683433010756988245264023278979171257383
9393665327679833614768150052656; 1.8203612657402207213774563922839672577355
14974121838803919890716313324647266063 0.4968091033398407272419482986701929
212781875193685168509869508953140844315696385; -0.0770905148321886540541087
```

```

523876336814258346951775208500402191859854482720920784 1.017598766563264849
060098898539206054763830809952064465251034057834928452758467]
[0.16539373471855043235269990616314167267561648071405024983323074603349333
38170379 -3.012724179940403476541315694134726655752203181759187696181593844
746228164546646; 2.85428469479678773960257067108243022564615065209268018511
3914115151343195573857 -1.0462039624694713406245650802505332985403318664480
38104804024755763734668972962; 2.266313493148240471976257897193309155759051
09007925423806327113920505007619545 3.6026629760390427059694591454084956647
12933175099558639218742509961079862076226; -1.38225289927695647331240946244
0305524783696590504834523227654507283541156933768 -0.0277340048781703367735
7467290835124580414026900497477772119014738313236894082272]
[0.93729565590191772386822937907849249927160003551614099575761950448138189
43330874 -3.150867381144996182004750209236630684158204794066582530662474788
18717481015792; 3.664460098368130854025896462345709559380591309280388130206
76783606788305457966 -0.433245272557746429190031899107521904504141721007496
6765392170486348240038260203; 1.9278353305372252613311284178556741305513578
41161844605750066192084743406786883 4.7077613564906059689244524472220425515
70102598317781549956496293646862232628021; -1.77978461190247321121473899601
5805096192365687274023095287964146660932730721375 -0.6320709452244916263568
030768271175565166180888262826559269062146386574813749507]

```

```
sol[1,3]
```

```

-0.275053974865865550518911557973824525679493409739538749452662408882393229
6471318

```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```

prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)

retcode: Success
Interpolation: Automatic order switching interpolation
t: 6-element Array{BigFloat,1}:
 0.0
 0.067107115562438418377697552621102545570540441832212517454262105079696189
49759242
 0.287686512296423947506690006844863599463955345609664823575147787622855867
5940776
 0.583118348391001237223137562711281971719124172681811699301835255327879129
7493794
 0.908644236705458802225160205670866220622479461008606364552372147659270654
6806487
 1.0
u: 6-element Array{Array{BigFloat,2},1}:
 [0.200754840726371952541740029118955135345458984375 0.77574051057756587823
632798972539603710174560546875; 0.89674334364338226244228735595243051648139
95361328125 0.9671172949562656384614456328563392162322998046875; 0.98052592
76780760036018591563333757221698760986328125 0.3303948714318076795848355686
7577135562896728515625; 0.1677649362171667046794709676760248839855194091796

```

```

875 0.070391167264269416392608036403544247150421142578125]
[0.13836592436844160367489127851113916678416026527029526355140858673807329
41241007 0.7728229813739593389306118953022683483641784084569804507430339699
882695271173377; 1.02733484649925107244485235774929366819580047260049886985
7308249851748568880979 1.06731565429613635487572454395633941827347436491014
6670930548999738326826543486; 0.9489493766100118783900414867981114598781255
891668659714322902078631061966671051 0.132186526396001050863829736215194773
5976338375870542007996833615079580662116519; 0.2695107459447620997748453716
88851423899247529608780853374039173225112567850078 0.2582877796952493502315
476023150140912742533800966009341819344611792211311349785]
[-0.2750539748658654899313378559291296284705986793802995902411424009457833
832601109 0.305387352130809343060642557167356574224650780560816481589523651
2164143390820049; 1.1178595289049776082018258826741252700246161210309610439
42679628014584997254817 0.7057055783332082566343028723716255795656896158142
665088553103942636145109180665; 1.07208851996608951131602444482730261458568
9542456058977572321435453223675588866 -0.2753515953994391249244732571356700
836252559279369040403857800732318800569262257; 0.40160411729150075300291604
31962070178186166895516082449090655329231546506637422 0.7988449318850210435
038454144001704570034874956649562373828384830294620391600754]
[-0.7847703693366268117340209397466439156080343828099543320468642427767410
707830353 -1.25801418559606860970797553528012902390378859586637511382936371
8158726036875872; 1.2321962098943288086663741827090796971274596838921114058
90900834640983791737064 -0.715417723683432609444608542194228307360172999343
0135844492780505542680546674509; 1.8203612657402204580664472118570278689817
18202784221045724886329597525672617088 0.4968091033398402070114742481154064
223454487856297059218880620575848238278690338; -0.0770905148321883978975346
9618969834196268571694189064000145289581652460037386879 1.01759876656326491
8957347887568284960250244750654869053503423114614299608779142]
[0.16539373471854949063273759595861178454360989988605843843163921745451003
08033735 -3.012724179940403094079223511049046002981465126109852951454658264
587588652834985; 2.85428469479678665305014713090025803185832044655639189955
158101877172436288295 -1.04620396246947194416427344306949458712220635510277
7015208808092626072806562005; 2.2663134931482407439569113216023810972571186
31461407067331275165524695806126371 3.6026629760390411059979649229894473272
05504217120429758181116305075852129700401; -1.38225289927695587504688940839
9586579079679433483263411507493121529576673868884 -0.0277340048781695584112
965231387219364208244053032074159318737809387866042673811]
[0.93729565590191745085831100267504022254411801334522701725378637208323945
65287039 -3.150867381144996182268169752252366014886891033316734330604404625
881618066161608; 3.66446009836813059115409234016169713168496659926410267812
9370724280026872461144 -0.4332452725577466787204600609486728737495399824367
111619158625626217368993514798; 1.92783533053722541479089770109675652377964
1892263121342582010620100911560026374 4.70776135649060563665151347493333676
9342329517188947115872874640332148469385822; -1.779784611902473096691807919
857679615333171108164504251701914110167906979845431 -0.63207094522449142183
88343066358528836589395323010878674009910269011525937558268]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)

```

```

tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

```

```
retcode: Success
```

```
Interpolation: Automatic order switching interpolation
```

```
t: 11-element Array{Float64,1}:
```

```

0.0
0.04077144418220069
0.09797709764732387
0.16362990401992172
0.240726389022935
0.34029638631605874
0.4648937736100185
0.6137018349351122
0.7787988400313253
0.9490115787984593
1.0

```

```
u: 11-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
```

```

[0.7896883235841079 0.8758256609227859; 0.5833739882166069 0.4868377261122
385; 0.07468513004306354 0.304186835584229; 0.724374236741433 0.59605244090
00487]
[0.6505800318720438 0.7640022115816141; 0.5582144852360046 0.5326217850716
795; -0.00967569863595065 0.19884722701311416; 0.9286526781551622 0.8296379
226578134]
[0.3756769761169915 0.5186400567498405; 0.42683310691591175 0.485894516724
80516; -0.06792049082502699 0.10663271837846786; 1.1978921924988644 1.14237
37364024698]
[-0.05231673049658053 0.1104917349828276; 0.15382395646550784 0.2866393663
772453; -0.025106198510803625 0.10700753436916774; 1.470779308514076 1.4673
195009726654]
[-0.6985712890715667 -0.5335724659877886; -0.29663654829678743 -0.11097755
773369194; 0.21054507733987743 0.293624432100265; 1.7227396507593182 1.7810
91564811748]
[-1.727477698531247 -1.5956935028197958; -0.9912204935568418 -0.7911935221
111379; 0.8692593093672187 0.9015308561704477; 1.9002065375394221 2.0349052
297360117]
[-3.210006709696324 -3.177500958116653; -1.802443211652017 -1.660864210594
3447; 2.329487132886209 2.3402479798652185; 1.8108858088146558 2.0260953866
84216]
[-4.955060440262974 -5.115636236851844; -2.175327004357237 -2.179290667651
5864; 4.998834783140636 5.0810744816139035; 1.1215648897015242 1.3931318749
124386]
[-6.20395187591461 -6.637629387710871; -0.9995599022573427 -1.177873925535
9167; 8.809311137921831 9.13084901201097; -0.5296334165138357 -0.2733579036
3937984]
[-5.716435661146319 -6.4192274339595565; 2.8630730835783207 2.622460845339
1243; 12.69573871328679 13.442158625471304; -3.2410596132734986 -3.11614375
62839547]
[-5.041512489387573 -5.805494137957469; 4.636997407095338 4.42355030945225
; 13.607572524377087 14.50854858009318; -4.225855877356378 -4.1666621846576
45]

```

```
sol[3]
```

```
4×2 StaticArrays.SArray{Tuple{4,2},Float64,2,8} with indices SOneTo(4)×SOne
To(2):
 0.375677  0.51864
 0.426833  0.485895
-0.0679205 0.106633
 1.19789   1.14237
```

0.3 Conclusion

These are the basic controls in DifferentialEquations.jl. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

0.4 Appendix

This tutorial is part of the DiffEqTutorials.jl repository, found at: <https://github.com/JuliaDiffEq/DiffEqTutorials>

To locally run this tutorial, do the following commands:

```
using DiffEqTutorials
DiffEqTutorials.weave_file("introduction","01-ode_introduction.jmd")
```

Computer Information:

```
Julia Version 1.4.2
Commit 44fa15b150* (2020-05-23 18:35 UTC)
Platform Info:
```

```
 OS: Linux (x86_64-pc-linux-gnu)
CPU: Intel(R) Xeon(R) CPU @ 2.30GHz
WORD_SIZE: 64
LIBM: libopenlibm
LLVM: libLLVM-8.0.1 (ORCJIT, haswell)
```

Environment:

```
JULIA_CUDA_MEMORY_LIMIT = 536870912
JULIA_DEPOT_PATH = /builds/JuliaGPU/DiffEqTutorials.jl/.julia
JULIA_PROJECT = @.
JULIA_NUM_THREADS = 4
```

Package Information:

```

Status `~/builds/JuliaGPU/DiffEqTutorials.jl/Project.toml`
[2169fc97-5a83-5252-b627-83903c6c433c] AlgebraicMultigrid 0.3.0
[7e558dbc-694d-5a72-987c-6f4ebed21442] ArbNumerics 1.0.5
[6e4b80f9-dd63-53aa-95a3-0cdb28fa8baf] BenchmarkTools 0.5.0
[be33ccc6-a3ff-5ff2-a52e-74243cff1e17] CUDAnative 3.1.0
[159f3aea-2a34-519c-b102-8c37f9878175] Cairo 1.0.3
[3a865a2d-5b23-5a0f-bc46-62713ec82fae] CuArrays 2.2.1
[55939f99-70c6-5e9b-8bb0-5071ed7d61fd] DecFP 0.4.10
[abce61dc-4473-55a0-ba07-351d65e31d42] Decimals 0.4.1
[ebbdde9d-f333-5424-9be2-dbf1e9acfb5e] DiffEqBayes 2.15.0
[eb300fae-53e8-50a0-950c-e21f52c2b7e0] DiffEqBiological 4.3.0
[459566f4-90b8-5000-8ac3-15dfb0a30def] DiffEqCallbacks 2.13.3
[f3b72e0c-5b89-59e1-b016-84e28bfd966d] DiffEqDevTools 2.21.0
[9fdde737-9c7f-55bf-ade8-46b3f136cc48] DiffEqOperators 4.10.0
[1130ab10-4a5a-5621-a13d-e4788d82bd4c] DiffEqParamEstim 1.14.1
[055956cb-9e8b-5191-98cc-73ae4a59e68a] DiffEqPhysics 3.2.0
[0c46a032-eb83-5123-abaf-570d42b7fbaa] DifferentialEquations 6.14.0
[31c24e10-a181-5473-b8eb-7969acd0382f] Distributions 0.23.4
[497a8b3b-efae-58df-a0af-a86822472b78] DoubleFloats 1.1.12
[f6369f11-7733-5829-9624-2563aa707210] ForwardDiff 0.10.10
[7073ff75-c697-5162-941a-fcdaad2a7d2a] IJulia 1.21.2
[23fbe1c1-3f47-55db-b15f-69d7ec21a316] Latexify 0.13.5
[c7f686f2-ff18-58e9-bc7b-31028e88f75d] MCMCChains 3.0.12
[eff96d63-e80a-5855-80a2-b1b0885c5ab7] Measurements 2.2.1
[961ee093-0014-501f-94e3-6117800e7a78] ModelingToolkit 3.1.1
[2774e3e8-f4cf-5e23-947b-6d7e65073b56] NLSolve 4.4.0
[429524aa-4258-5aef-a3af-852621145aeb] Optim 0.21.0
[1dea7af3-3e70-54e6-95c3-0bf5283fa5ed] OrdinaryDiffEq 5.41.0
[65888b18-ceab-5e60-b2b9-181511a3b968] ParameterizedFunctions 5.3.0
[91a5bcdd-55d7-5caf-9e0b-520d859cae80] Plots 1.4.0
[d330b81b-6aea-500a-939a-2ce795aea3ee] PyPlot 2.9.0
[731186ca-8d62-57ce-b412-fbd966d074cd] RecursiveArrayTools 2.4.4
[47a9eef4-7e08-11e9-0b38-333d64bd3804] SparseDiffTools 1.8.0
[684fba80-ace3-11e9-3d08-3bc7ed6f96df] SparsityDetection 0.3.2
[90137ffa-7385-5640-81b9-e52037218182] StaticArrays 0.12.3
[f3b207a7-027a-5e70-b257-86293d7955fd] StatsPlots 0.14.6
[c3572dad-4567-51f8-b174-8c6c989267f4] Sundials 4.2.2
[1986cc42-f94f-5a68-af5c-568840ba703d] Unitful 1.2.1
[44d3d7a6-8a23-5bf8-98c5-b353f8df5ec9] Weave 0.10.2
[b77e0a4c-d291-57a0-90e8-8db25a27a240] InteractiveUtils
[37e2e46d-f89d-539d-b4ee-838fcccc9c8e] LinearAlgebra
[44cfe95a-1eb2-52ea-b672-e2afdf69b78f] Pkg

```